

Solution Abstract Problem 2: Almost Famous

Cindy Lamm

Tuesday 20th January, 2015

1 Approach

I identified question 1 as classification problem (supervised or unsupervised) which I wanted to solve using Spark (after the performance problems I encountered for the classification in R with 7.3 million observations).

I opted for the following approach to solve the questions:

- analyse file structure
- explore data
- prepare data for modeling
- estimate and validate models for spam classification
- prepare the web data for prediction
- classify entries in web data as spam
- remove spam entries from web data
- calculate required metrics

1.1 Analyse file structure

I had a short look into the log data in json format via `jq`. Once I saw that the data is very regular (same number of keys per row, same names for keys) I converted the data to csv¹ using `json2csv` (which is pre-installed in `the data science toolbox`) to load it into R for exploratory analysis.

1.2 Explore data

To get an impression of the web and spam data I had a look at overall summary statistics but also at summary statistics aggregated over visits and visitors. Since I am not that familiar with web analytics I assume that a uid uniquely identifies a visitor which is a synonym for user.

From manual inspection it seems that visits that are generated via the bot network only contain the action `adclick` after they landed on the page. However if I would classify the web data according to this simple rule, then the click through rate for ads on the page would be zero - which seems too extreme.

¹see `problem2_famous/data/convert2csv.sh`

1.3 Prepare data for modeling

In order to run a classification/clustering algorithm in Spark I need to code the categorical variables (all but `tstamp`) with integer levels (which means that I implicitly assume that it is okay if an algorithm treats these variables as continuous features, because the resulting model is still good enough). I converted the `tstamp` to epoch values (i.e. to seconds since 1970-01-01 00:00:00 GMT) to have a true continuous variable without losing the time information.

For the supervised classification I added a variable `is_spam` with value 1 to the spam data and with value 0 to the web data (knowing that with this assumption some web data is falsely labeled). I pooled the spam data and different amounts of the web data into a training data sets - which I used as well for testing (see section 1.5).

1.4 Estimate and validate models for spam classification

I tried the different classification and cluster algorithms: I started with unsupervised clustering using k-means (with $k = 2$) on a training sample with all spam log data and (the top) 200.000 observations of the web log data² I had the hope that by using $k = 2$ all spam data would end up in one cluster and all web data in the other, so that when I predict the complete web data on this clustering all so far unknown spam entries go into the spam-cluster. However when I used the estimated cluster center to classify the training data, the result was not as good as I would have expected:

TODO: add venn diagram here or error/accuracy values

1.5 Current Limitations Of Spam Classification

Currently I used the same data for training, validation and prediction, which decreases the generalization performance of the estimated model. It would be better to split the available spam and web log data into independent sets for training and validation. For classification/clustering one could use the time difference between the actions and/or the total time spent per visit as input. I also deem worth trying to cluster/classify on aggregated visit and/or user data. Given that I relabeled the categorical variables with (rather small) integer labels it might be necessary to scale the epoch values, because they otherwise bias the differences. It might also be better to use 1-of-n encoding for the categorical variables instead of just giving them numerical labels.

1.6 Compute Required Metrics

Calculate Overall Click-Trough Rate

I assume the following formula for the overall click-through rate (CTR), that was described as "ad clicks per visit"

$$CTR = \frac{\#\{action = adclick\}}{\#\{unique\ visit\ ids\}}$$

Before having removed the spam visits from the web log data I counted 103896 ad clicks and 1482602 unique visit ids, which leads to a CTR of 7.01% - which is unusually high if we believe in wikipedia³.

After having removed the spam visits from the web log data I counted `xx` ad clicks and `xx` unique visit ids, which leads to a CTR of `xx%`.

Calculate Mean Orders Per Visit Within Campaign-Query Combinations

Since the action "order" is only listed in log entries that don't have query and campaign information the first step I did was to aggregate the web log per visit (using Spark `reduceByKey` method on the key (`visit_id,uid`)), putting all actions in a list. I saved the aggregated visits data back to a log file, converted it to csv and work on that in R, where I feel more comfortable computing statistics per groups compared to map/reduce paradigm necessary in Spark.

²Using random 200.000 observations from the web log data didn't seem to be a good idea given the timely ordered structure.

³http://en.wikipedia.org/wiki/Click-through_rate

Possible Improvements

I wrote a separate Spark Python script for each of the metrics where I loaded the data from the file system in each script. This could be improved by calculating all metrics within one Spark session for which I would load the data only once from the file system.

2 Used Software

In general I developed on a Macbook Pro with a 2.3 GHz Intel Core i7 Processor and 16GB Memory.

- git
- jq
- json2csv (pre-installed in the data science toolbox)
- LaTeX
- R in RStudio with packages knitr, plyr
- shell command line
- Spark with Python API and MLlib package

3 Time Spent

-

Summing up I spend **xx hours** on this problem.