

Almost Famous: Analyse campaign query combinations

Cindy Lamm

12:35, Wednesday 21st January, 2015

Load variable names and types:

```
nameTypeDataFile <- "../data/raw_variables.csv"
variableNames <- read.csv(nameTypeDataFile, header=TRUE, stringsAsFactors=FALSE)
variableNames

##      name      type
## 1  visit_id  factor
## 2      uid   factor
## 3  campaign  factor
## 4    tstamp character
## 5 experiments  factor
## 6     action  factor
## 7      query  factor

factorIdx <- which(variableNames$type=="factor")
factorNames <- variableNames$name[factorIdx]
```

Read the per visit aggregated web log data:

```
visitFile <- "../data/web_visits.csv"
visitData <- read.csv(visitFile, stringsAsFactors=FALSE, col.names=variableNames$name,
                      colClasses=variableNames$type, na.strings=c("NA",""))
visitData$tstamp <- as.POSIXct(visitData$tstamp)
str(visitData)

## 'data.frame': 1482602 obs. of 7 variables:
## $ visit_id : Factor w/ 1482602 levels "10000024498",...: 1252062 128641 583195 349394 830165 690964 ...
## $ uid      : Factor w/ 1064214 levels "100000493","100000682",...: 858988 92339 95584 929716 656934 ...
## $ campaign : Factor w/ 10 levels "103","127","14",...: 7 1 7 1 4 8 1 1 1 2 ...
## $ tstamp   : POSIXct, format: "2014-09-18 05:43:18" "2014-09-16 21:24:08" ...
## $ experiments: Factor w/ 4 levels "[1 3]","[1 4]",...: 2 1 4 1 3 2 1 3 2 1 ...
## $ action    : Factor w/ 8 levels "[landed adclick adclick adclick]",...: 8 8 8 8 8 8 8 8 8 8 ...
## $ query     : Factor w/ 5 levels "advanced analytics",...: 1 5 1 5 3 5 5 5 5 4 ...
```

```
summary(visitData)

##          visit_id          uid          campaign          tstamp
## 10000024498:      1 102486699:      7 558      :324872 Min.      :2014-09-15 00:00:01
## 10000032484:      1 123618732:      7 103      :324027 1st Qu.:2014-09-18 16:32:04
## 10000079220:      1 143588980:      7 59      :232002 Median :2014-09-22 16:55:36
## 10000092303:      1 159226004:      7 31      :231685 Mean   :2014-09-22 20:33:11
## 10000132469:      1 168873739:      7 127      : 92681 3rd Qu.:2014-09-26 19:41:15
## 10000206890:      1 171898393:      7 94      : 92436 Max.   :2014-09-30 23:53:20
## (Other)      :1482596 (Other)      :1482560 (Other):184899
## experiments          action
## [1 3]:370018 landed          :1291256
## [1 4]:371852 [landed signup]      : 84889
## [2 3]:370082 [landed order]        : 43930
## [2 4]:370650 [landed adclick]       : 28233
##           [landed adclick adclick adclick]: 14956
##           [landed adclick adclick]        : 14875
##           (Other)                        : 4463
##           query
## advanced analytics      :463687
## building predictive models: 92454
## data science            : 92445
## data science training    :185117
## predictive modeling      :648899
##
##
```

What are the actions per visit??

```
table(visitData$action)

##
## [landed adclick adclick adclick]          [landed adclick adclick]
##                               14956                14875
##           [landed adclick]                [landed order]
##                               28233                43930
##           [landed signup adclick]          [landed signup order]
##                               1045                3418
##           [landed signup]                  landed
##                               84889                1291256
```

Look at visits with orders:

```
orderLevels <- names(unlist(sapply(levels(visitData$action), FUN=grep, pattern="order")))
orderLevels

## [1] "[landed order]"      "[landed signup order]"

isOrderIdx <- sort(union(which(visitData$action == orderLevels[1]),
                             which(visitData$action == orderLevels[2])))
totalOrders <- length(isOrderIdx)
```

I conclude from the factor levels for action that there is at most 1 order per visit and overall 47348 orders. I cross check with a simple grep on the command line on the unaggregated web data which gives us the same result:

```
$ grep -o order web.log | wc -l
$ 47348
```

Add the number of orders per visit as variable to the data frame:

```
nbOrder <- rep(0, nrow(visitData))
nbOrder[isOrderIdx] <- 1
visitData$nb_orders <- nbOrder
```

```
prop.table(table(visitData$nb_orders))
```

```
##
##          0          1
## 0.96806425 0.03193575
```

There are 96.8064255% of visits that don't have an order and only 3.1935745% that do.
How many orders are there per campaign-query combination?

```
combinations <- expand.grid(queries=levels(visitData$query), campaigns=levels(visitData$campaign))
length(combinations)
```

```
## [1] 2
```

```
webAggCampaignQuery <- aggregatePerCQ(visitData)
o <- order(webAggCampaignQuery$mean_orders, decreasing=TRUE)
webAggCampaignQuery[o,]
```

```
##      campaign               query nb_visits nb_uids total_orders mean_orders
## 3      14 building predictive models    46252   45738         3065  0.06626740
## 2      127      data science training    92681   90761         5101  0.05503825
## 10      94      data science training    92436   90510         5071  0.05485958
## 5      203 building predictive models    46202   45711         2511  0.05434830
## 8      558      predictive modeling    324872  301174        10628  0.03271442
## 1      103      predictive modeling    324027  300394        10425  0.03217324
## 4       17           data science     46308   45814         1052  0.02271746
## 9       59      advanced analytics    232002  219833         4396  0.01894811
## 6       23           data science     46137   45630          858  0.01859679
## 7       31      advanced analytics    231685  219610         4241  0.01830503
##      sd_orders
## 3  0.2487516
## 2  0.2280561
## 10 0.2277072
## 5  0.2267061
## 8  0.1778884
## 1  0.1764603
## 4  0.1490029
## 9  0.1363421
## 6  0.1350975
## 7  0.1340523
```

Write the result into csv file:

```
write.csv(webAggCampaignQuery, file="../q3_campaign_query_combi/webAggCampaignQuery.csv",  
          row.names=FALSE, quote=FALSE)
```

I use a Python script to put the result in the required json format (because even if I would use the package jsonlite to format the result in R, the `sink()` method, which I would need to write it to the file system, does not work in combination with knitr.)